

(COVER (mandatory item – NBR14724, item 4.1.1))

Daniel Quintão Dávila

[An Exploration of the Evolving Landscape of Artificial Intelligence Agents]
(required item; font size 12; bold style) : [Capabilities, Challenges, and Future Prospects], if any **(font size 12; normal style)**

SÃO PAULO

2025

(COVER PAGE (mandatory item – NBR14724, item 4.2.1.1))

Daniel Quintão Dávila

**[An Exploration of the Evolving Landscape of Artificial Intelligence Agents]
(required item; font size 12; bold style) : [Capabilities, Challenges, and Future
Prospects], if any (font size 12; normal style)**

Final Course Project submitted to the
Institute of Technology and Leadership
(INTELI), to obtain a bachelor's degree in
Computer Engineering. **(font size 12;
normal style; single spacing; NBR
14724, item 5)**

Advisor: Prof. Rafael Matsuyama

Coadvisor: Prof. [professor's name]

SÃO PAULO
2025

(Catalog Card (mandatory item – NBR14724, item 4.2.1.1.2))

Dávila, Daniel

An Exploration of the Evolving Landscape of Artificial Intelligence Agents: Capabilities, Challenges, and Future Prospects / Daniel Quintão Dávila; Prof. Rafael Matsuyama — São Paulo, 2025.

Nº de páginas :

Trabalho de Conclusão De Curso (Graduação) — Curso de Engenharia de Hardware
/ Instituto de Tecnologia e Liderança
Bibliografia

Cataloging in Publication
Library and Documentation Service
Instituto de Tecnologia e Liderança (INTELI)
Data entered by the author.

(Cataloging record with international cataloging data, according to NBR 14724. The record will be completed later, after approval and before the final version is deposited. The completion of the cataloging record is the responsibility of the institution's library.)

Sobrenome, Nome

Título do trabalho: subtítulo / Nome Sobrenome do autor; Nome e Sobrenome do orientador. — São Paulo, 2025.

nº de páginas : il.

Trabalho de Conclusão de Curso (Graduação) – Curso de [Ciência da Computação] [Engenharia de Software] [Engenharia de Hardware] [Sistema de Informação] / Instituto de Tecnologia e Liderança.

Bibliografia

1. [Assunto A]. 2. [Assunto B]. 3. [Assunto C].

CDD. 23. ed.

Resumo (**mandatory item – NBR 14724, item 4.2.1.7**)

Dávila, Daniel. **An Exploration of the Evolving Landscape of Artificial Intelligence Agents.** 2025. nº de folhas. TCC (Graduação) – Curso Engenharia de Hardware, Instituto de Tecnologia e Liderança, São Paulo, 2025. **(The reference is an optional item. NBR 6028, item 4.1.6)**

A Inteligência Artificial (IA) tornou-se uma tecnologia fundamental da era moderna, impulsionando o desenvolvimento de sistemas autônomos sofisticados. Esta tese apresenta uma revisão abrangente e introdutória desses sistemas, conhecidos como agentes de IA. O objetivo principal deste trabalho é sintetizar a literatura consolidada a fim de explorar os conceitos centrais, as arquiteturas, o contexto histórico e as implicações no mundo real da tecnologia de agentes de IA.

A metodologia deste trabalho consiste em uma revisão estruturada. Inicia-se com uma revisão da literatura relacionada e um levantamento da evolução histórica dos sistemas autônomos. Em seguida, são estabelecidas as definições formais de agentes, racionalidade e do framework PEAS (Medida de Desempenho, Ambiente, Atuadores, Sensores). A tese categoriza sistematicamente as principais arquiteturas de agentes, detalhando a progressão de Agentes de Reflexo Simples para sistemas mais complexos baseados em utilidade. Esses conceitos teóricos são então aplicados por meio de estudos de caso aprofundados de um agente de xadrez, um aspirador robô e um veículo autônomo.

Por fim, a tese conclui examinando os desafios técnicos significativos e as profundas considerações éticas que acompanham essa tecnologia, oferecendo uma análise ampliada de questões como viés algorítmico, privacidade do usuário, substituição de postos de trabalho e responsabilização legal. O trabalho conclui que, embora agentes de IA ofereçam potencial transformador, seu desenvolvimento exige governança cuidadosa.

Palavras-Chave: [Inteligência Artificial]; [Agentes de IA]; [PEAS]; [Sistemas Autônomos]. **(Required item – NBR 6028, item 4.1.7; keywords must be written with lowercase initials, except for proper nouns, and separated by semicolons and ending with a period)**

Abstract (mandatory item – NBR 14724, item 4.2.1.8)

Dávila, Daniel. **An Exploration of the Evolving Landscape of Artificial Intelligence Agents.** 2025. n° of pages. Final course project (Bachelor) – Hardware Engineering, Institute of Technology and Leadership, São Paulo, 2025. **(The reference is an optional item. NBR 6028, item 4.1.6)**

Artificial Intelligence (AI) has become a foundational technology of the modern era, leading to the development of sophisticated autonomous systems. This thesis provides a comprehensive, foundational review of these systems, known as AI Agent. The primary objective of this paper is to synthesize established literature in order to explore the core concepts, architectures, historical context, and real-world implications of AI agent technology.

The methodology of this paper is a structured review. It begins with a review of related literature and a survey of the historical evolution of autonomous systems. It then establishes the formal definitions of agents, rationality, and the PEAS (Performance Measure, Environment, Actuators, Sensors) framework. The thesis systematically categorizes primary agent architectures, detailing the progression from Simple Reflex Agents to more complex Utility-Based systems. These theoretical concepts are then applied through in-depth case studies of a chess agent, a robotic vacuum, and an autonomous vehicle.

Finally, the thesis concludes by examining the significant technical challenges and profound ethical considerations that accompany this technology, providing expanded analysis on issues of algorithmic bias, user privacy, job displacement, and legal accountability. The paper finds that while AI agents offer transformative potential, their development requires careful governance.

Key words: [Artificial Intelligence]; [AI Agents]; [PEAS]; [Autonomous Systems]. **(Required item – NBR 6028, item 4.1.7; keywords must be written with lowercase initials, except for proper nouns, and separated by semicolons and ending with a period)**

Summary (mandatory item – NBR14724, item 4.2.1.13; NBR 6027)

Contents

Abstract i

1 Introduction

1.1 Background: The Age of Artificial Intelligence

1.2 The Emergence of Autonomous Systems

1.3 Thesis Objective

1.4 Definitions of Key Terms

1.5 Structure of the Thesis

2 Review of Related Literature

2.1 Introduction

2.2 Foundational and Theoretical Literature

2.3 Industrial and Application-Focused Literature

2.4 Ethical and Societal Critiques

2.5 Summary

3 A Historical Perspective on Autonomous Systems

3.1 Introduction

3.2 The Birth of Artificial Intelligence (The 1950s)

3.3 Early Successes and Fundamental Limitations (The 1960s)

- 3.4 The First AI Winter (The 1970s)
- 3.5 The Rise of Expert Systems (The 1980s)
- 3.6 The Second AI Winter and the Shift to Agents (Late 1980s
– Mid 1990s)
- 3.7 The Modern Renaissance: Big Data and Deep Learning (Present)

4 Foundations of Artificial Intelligence Agents

- 4.1 Defining the "Agent"
- 4.2 From Agents to Rational Agents
- 4.3 The PEAS Framework
 - 4.3.1 P — Performance Measures
 - 4.3.2 E — Environment
 - 4.3.3 A — Actuators
 - 4.3.4 S — Sensors
- 4.4 The Importance of a PEAS-First Approach to the Modelling of AI Agents

5 Architectures of AI Agents

- 5.1 Simple Reflex Agents
- 5.2 Model-Based Reflex Agents
- 5.3 Goal-Based Agents
- 5.4 Utility-Based Agents
- 5.5 Summary of Architectures

6 In-Depth Case Studies of Agent Implementations

- 6.1 Introduction to Case Studies
- 6.2 Case Study 1: The Chess-Playing Agent
 - 6.2.1 PEAS Analysis
 - 6.2.2 Environmental Classification
 - 6.2.3 Architectural Analysis
- 6.3 Case Study 2: The Autonomous Robotic Vacuum Cleaner
 - 6.3.1 PEAS Analysis
 - 6.3.2 Environmental Classification
 - 6.3.3 Architectural Analysis
- 6.4 Case Study 3: The Self-Driving Car

6.4.1 PEAS Analysis

6.4.2 Environmental Classification

6.4.3 Architectural Analysis

7 Applications and Societal Impact

7.1 Real-World Applications

7.2 Technical Challenges

7.3 Ethical Considerations in Autonomous Systems

7.3.1 Bias, Fairness, and Algorithmic Discrimination

7.3.2 Privacy, Consent, and the Surveillance State

7.3.3 Job Displacement and Economic Disruption

7.3.4 The Accountability Gap

8 Conclusion and Future Outlook

8.1 Summary of Thesis

8.2 The Future of AI Agents

8.3 Final Word

Bibliography

1. Introduction (Titles and numeric indicative see NBR 14724, item 5.2.2; sequential numbering see NBR6024)

1.1 Background: The Age of Artificial Intelligence

We live in a time of unprecedented technological change, an era often referred to as the “Digital Revolution” or the “Fourth Industrial Revolution.” This period is characterized by the convergence of digital, physical, and biological systems, fundamentally altering how we live, work, and interact. At the very heart of this transformation is the field of computer science known as Artificial Intelligence (AI). Once a concept relegated to the realm of speculative science fiction, AI has transitioned into a tangible and powerful reality, becoming one of the most critical and discussed technologies of the twenty-first century.

The broad field of Artificial Intelligence is concerned with a monumental task: the creation of non-biological, or “artificial,” systems that can simulate or replicate tasks that typically require human intelligence. This encompasses a wide spectrum of capabilities, from basic pattern recognition and logical reasoning to more complex functions like language comprehension, decision-making, and visual perception (Frank, 2019). As a discipline, AI is not a single technology but rather an umbrella term that includes many sub-fields, most notably machine learning, deep learning, natural language processing, and computer vision.

The recent acceleration of AI’s development and adoption can be attributed to several key factors. First, the proliferation of the internet and smart devices has generated an almost incomprehensible amount of data (often termed “Big Data”), which serves as the raw material for training modern AI systems. Second, significant advancements in computing power, particularly the development of specialized hardware like Graphics Processing Units (GPUs), have made it possible to process this data at a scale and speed previously unattainable. Together, these developments have fueled the “AI boom,” moving the technology from academic laboratories into practical, everyday applications that shape the modern world. It is within this expansive and revolutionary context that we begin our specific inquiry into AI agents.

1.2 The Emergence of Autonomous Systems

The practical application of Artificial Intelligence, as discussed in the previous section, is not limited to passive data analysis or information retrieval. The field is witnessing a significant and transformative paradigm shift, moving from systems that merely process information to systems that can act upon it. This evolution has led to the emergence of autonomous systems, which are the central focus of this thesis.

The concept of autonomy, in this context, refers to a system's capacity to operate independently and make decisions without direct, moment-to-moment human intervention. These autonomous systems are the digital or physical embodiments of AI principles, and they are more formally known as agents. An agent is, in its most basic sense, an entity that perceives its surroundings and takes actions to achieve a specific objective.

This represents a fundamental departure from traditional computing models. Previously, software was largely passive; it functioned as a tool that required a human operator to provide explicit instructions and interpret the results. In contrast, the autonomous systems now emerging are designed to be active participants in their environments. They are becoming collaborators, assistants, and delegates for human tasks in both the digital realm (such as software bots) and the physical world (such as robots and autonomous vehicles). This transition from passive tools to active, goal-oriented agents is a major step in the evolution of technology, and it is the reason this topic demands specific study.

1.3 Thesis Objective

The primary objective of this thesis is to provide a comprehensive, foundational overview of the field of Artificial Intelligence agents. This paper does not seek to introduce novel experimental research or groundbreaking analysis. Instead, its purpose is to synthesize and present the established, core concepts that define this critical area of computer science in a clear and structured manner. It serves as a comprehensive review of the existing literature and foundational principles.

To this end, this thesis will:

- Review the existing academic and industrial literature regarding AI agents
- Explore the historical context and foundational definitions of rational, autonomous agents
- Categorize the primary architectures of AI agents, from simple reflex models to complex utility-based systems
- Analyze concrete examples of agents through detailed case studies
- Discuss real-world applications, technical challenges, and provide an in-depth examination of the ethical considerations of autonomous technology

The central goal is to provide the reader with a clear and accessible understanding of what AI agents are, how they are structured to function, and the role they are beginning to play in the modern technological landscape.

1.4 Definitions of Key Terms

To ensure clarity throughout this thesis, the following fundamental terms are defined as they are used in the context of Artificial Intelligence research. These definitions will be expanded upon in subsequent chapters, drawing heavily from standard texts in the field (Russell and Norvig, 2020).

- Artificial Intelligence (AI): A sub-field of computer science aimed at creating systems capable of performing tasks that typically require human intelligence, such as learning, reasoning, and perception.
- Agent: In the broadest sense, any entity that perceives its environment through sensors and acts upon that environment through actuators.
- Rational Agent: An agent that acts so as to maximize the expected value of a performance measure, given the percept sequence it has observed and its built-in knowledge base. Rationality is concerned with the quality of decisions, not the thought process behind them.
- Environment: The external problem space in which an agent operates. The properties of the environment (such as observability and stochasticity) fundamentally dictate the required complexity of the agent's architecture.

- Percept: The input that an agent's sensors provide at any given instant regarding the state of the environment.
- Actuator: The mechanism by which an agent performs an action to affect change within its environment.
- Architecture: The internal structure of the agent; the specific computing machinery and software logic that maps percepts to actions.

1.5 Structure of the Thesis

To methodically achieve the stated objective, this thesis is organized into eight distinct chapters.

- Chapter 1: Introduction

Introduces the topic, defines key terms, and outlines the thesis objectives.
- Chapter 2: Review of Related Literature

Provides a survey of the key academic texts, industrial reports, and societal critiques that form the basis of current understanding of AI agents.
- Chapter 3: A Historical Perspective on Autonomous Systems

Reviews the evolution of AI research that led to the current agent paradigm.
- Chapter 4: Foundations of Artificial Intelligence Agents

Establishes the core theoretical definitions, including the concept of rationality, the PEAS framework, and environmental classification.
- Chapter 5: Architectures of AI Agents

Details the internal structures of agents, from simple reflex models to complex utility-based architectures.
- Chapter 6: In-Depth Case Studies of Agent Implementations

Applies the theoretical concepts from previous chapters to analyze real-world examples.

- Chapter 7: Applications and Societal Impact

Reviews current deployments of agent technology, outlines technical challenges, and examines key ethical considerations.

- Chapter 8: Conclusion and Future Outlook

Summarizes the key findings and offers a perspective on the future direction of the field.

2. Review of Related Literature

2.1 Introduction

The study of Artificial Intelligence Agents is a multidisciplinary endeavor that draws upon computer science, mathematics, philosophy, economics, and ethics. As such, the literature surrounding this topic is vast and varied, ranging from highly technical algorithmic analyses to broad societal critiques. This chapter provides a review of the key academic and industrial literature that forms the foundation for the subsequent analysis in this thesis. The literature can be broadly categorized into three areas: foundational and theoretical texts, industrial and application-focused reports, and ethical and societal critiques.

2.2 Foundational and Theoretical Literature

The theoretical bedrock for modern AI agent design is established in seminal works that define the scope and goals of the field. The most ubiquitous text in this domain, serving as the primary reference for the core definitions used in this thesis, is *Artificial Intelligence: A Modern Approach* by Stuart Russell and Peter Norvig (Russell & Norvig, 2020). Now in its fourth edition, this text is widely considered the standard in the field. Russell and Norvig are credited with popularizing the agent-based approach to AI, shifting the pedagogical focus from isolated algorithms to the concept of rational agents that perceive and act within an environment. Their formalization of the PEAS framework (Performance, Environment, Actuators, Sensors) and their taxonomy of agent architectures (reflex, goal-based, utility-based) provide the

essential vocabulary used by contemporary AI researchers and engineers (Russell & Norvig, 2020).

The concept of the agent itself is further explored in Michael Wooldridge's *An Introduction to Multi-Agent Systems* (Wooldridge, 2009). While Russell and Norvig focus heavily on the single rational agent, Wooldridge expands this view to consider the complexities that arise when multiple autonomous agents interact. His work is foundational for understanding competitive and cooperative agent environments, a topic that is increasingly relevant in domains such as autonomous driving and decentralized finance.

The historical roots of these theoretical concepts must also be acknowledged. Alan Turing's seminal 1950 paper, "Computing Machinery and Intelligence," published in *Mind*, laid the philosophical groundwork for the entire discipline (Turing, 1950). By proposing the "Imitation Game," now commonly known as the Turing Test, Turing shifted the discourse from abstract notions of thinking to observable behavior. This behavioral perspective is the direct ancestor of the modern definition of an AI agent, which is judged solely by the rationality of its actions rather than any internal cognitive state (Turing, 1950). Collectively, these foundational texts provide the rigorous framework necessary to define, design, and evaluate autonomous systems.

2.3 Industrial and Application-Focused Literature

Moving beyond theory, a significant body of literature focuses on the practical application, commercialization, and current state of AI agent technology in industry. These sources, often produced by technology research firms and business analysts, provide crucial insight into how theoretical constructs are deployed in real-world environments.

Gartner, Inc., a leading research and advisory company, regularly publishes reports defining strategic technology trends. Its 2023 report, *Top Strategic Technology Trends 2024: AI Trust, Risk and Security Management*, highlights the transition of AI from passive tools to active agents in enterprise environments (Gartner, 2023). This literature emphasizes the practical challenges of deployment, particularly in relation to security, reliability, and the management of autonomous decision-making processes in business-critical applications.

Similarly, business-oriented publications such as *Forbes* and *Business Horizons* frequently examine the immediate impact of AI agents on consumer markets and the workforce. Kaplan and Haenlein, writing in *Business Horizons*, provide a framework for understanding how AI is interpreted by end users, using ubiquitous examples such as Apple's Siri to illustrate the gap between public perception and technical reality (Kaplan & Haenlein, 2019). Bernard Marr, writing in *Forbes*, surveys current

use cases, demonstrating the pervasive role of utility-based agents in areas such as recommendation systems and logistics (Marr, 2020).

This body of industrial literature confirms that the agent architectures defined by Russell and Norvig are not merely academic abstractions but serve as functional blueprints for large-scale technological infrastructure. At the same time, these sources highlight the persistent gap between theoretical optimality and the complexities of real-world implementation.

2.4 Ethical and Societal Critiques

As AI agents have transitioned from controlled research environments to widespread real-world deployment, a substantial body of literature has emerged that examines their ethical implications and societal consequences. This literature challenges the purely technical view of agents as neutral optimization systems.

Meredith Broussard's *Artificial Unintelligence: How Computers Misunderstand the World* provides a critical counterpoint to the optimism often found in industrial reports (Broussard, 2018). Broussard argues that data-driven agents frequently reinforce existing societal biases and perform poorly in complex, context-dependent human situations. Her work highlights the limitations of the rational agent model when applied to inherently irrational social structures.

Luciano Floridi and Josh Cowls, writing in the *Harvard Data Science Review*, attempt to provide a constructive response to these challenges by proposing a unified framework of ethical principles for AI in society (Floridi & Cowls, 2019). They argue that the development of autonomous agents must be guided by principles such as beneficence, non-maleficence, autonomy, justice, and explicability. This perspective emphasizes that the design of utility functions is not merely a technical exercise but also a fundamentally moral one.

Ben Shneiderman further contributes to this discourse through his advocacy of a Human-Centered AI approach (Shneiderman, 2020). He argues that increasing levels of agent autonomy should be carefully balanced with mechanisms for human control and oversight, particularly in safety-critical domains. Collectively, this literature suggests that the successful deployment of AI agents depends not only on technical performance but also on trust, fairness, and accountability.

2.5 Summary

The literature on AI agents reflects a convergence of rigorous theoretical foundations, optimistic industrial applications, and critical ethical analysis. Foundational texts establish the core definitions and architectural models. Industrial literature demonstrates the real-world relevance and economic significance of these concepts. Ethical critiques provide necessary context and caution regarding their societal impact. This thesis draws upon all three strands of literature to provide a balanced and comprehensive examination of Artificial Intelligence agents.

3. A Historical Perspective on Autonomous Systems

3.1 Introduction

To fully understand the current capabilities, limitations, and future trajectory of modern Artificial Intelligence (AI) agents, it is essential to first examine the historical context from which this technology emerged. The concept of an autonomous, rational agent did not appear suddenly. Rather, it is the result of decades of iterative research, theoretical shifts, periods of intense optimism, and subsequent periods of disillusionment.

The history of AI is characterized by a continuous oscillation between the desire to create *general intelligence*—systems that can think and act like humans across any domain—and the practical reality of creating specialized systems that can perform specific tasks in restricted environments. The modern AI agent represents a convergence of these historical trends, leveraging advancements in computational power and algorithmic design to move closer to the goal of autonomous operation in complex, real-world environments. This chapter provides a chronological overview of the key developments that laid the conceptual groundwork for the agent-based systems we see today.

3.2 The Birth of Artificial Intelligence (The 1950s)

While the idea of artificial beings exists in ancient mythology and fiction, the formal scientific pursuit of Artificial Intelligence began in the mid-20th century, born out of the convergence of new theories in computation, logic, and cybernetics.

A seminal moment in the early history of the field was the publication of Alan Turing's 1950 paper, *Computing Machinery and Intelligence*. In this paper, Turing proposed the famous *Imitation Game*, now known as the Turing Test, as a criterion for intelligence. Crucially, Turing shifted the question from the philosophical "Can machines think?" to the operational "Can machines act indistinguishably from a human?" This focus on *acting*—on observable behavior—is a foundational concept

for the later development of AI agents, whose success is defined by their ability to act rationally within an environment.

The field officially received its name in 1956 at the Dartmouth Summer Research Project on Artificial Intelligence. Organized by John McCarthy, Marvin Minsky, Nathaniel Rochester, and Claude Shannon, this workshop brought together researchers interested in simulating various facets of human intelligence. The workshop proposal stated the ambitious hypothesis that “every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it.” This period was marked by immense optimism, with many researchers believing that human-level intelligence would be achieved within a generation.

3.3 Early Successes and Fundamental Limitations (The 1960s)

Following the Dartmouth workshop, the 1960s saw significant funding directed toward AI research. Researchers focused on developing programs that could solve problems in well-defined, controlled domains.

A key methodological approach during this period was the use of *microworlds*. These were simplified, artificial environments in which the rules were explicit and the state of the world was fully observable. A classic example was the *blocks world*, where a robotic arm was tasked with stacking blocks in a specific configuration. Programs developed at MIT, such as SHRDLU, demonstrated impressive capabilities within these restricted settings, including the ability to understand natural-language commands and manipulate objects accordingly.

From the perspective of agent theory, these early systems can be seen as precursors to modern agents. They perceived an environment, processed information, and executed actions to achieve a goal. However, their limitations were severe. They operated in environments that were static, deterministic, and fully observable—conditions that differ radically from the complexity of the real world. Moreover, they relied heavily on general-purpose search algorithms that did not scale, succumbing to combinatorial explosion when applied to larger or less constrained problems.

3.4 The First AI Winter (The 1970s)

By the early 1970s, the gap between the ambitious promises of early AI researchers and the practical performance of their systems had become impossible to ignore. Techniques that worked well in constrained microworlds failed when confronted with real-world complexity.

Two major factors contributed to the period of declining funding and confidence known as the *First AI Winter*. First, critical government evaluations—most notably the

1966 ALPAC report—highlighted limited progress in areas such as machine translation, leading to substantial funding cuts. Second, in 1969, Marvin Minsky and Seymour Papert published *Perceptrons*, which rigorously demonstrated theoretical limitations of early neural network models. Although the work was mathematically sound, it was widely interpreted as evidence that connectionist approaches were fundamentally flawed, causing neural network research funding to collapse for more than a decade.

3.5 The Rise of Expert Systems (The 1980s)

The 1980s marked a resurgence of interest in AI, driven largely by the commercial success of *expert systems*. This paradigm shifted focus away from general reasoning mechanisms toward systems densely packed with domain-specific knowledge.

An expert system is designed to replicate the decision-making abilities of a human expert within a narrowly defined field. Such systems typically consisted of two core components: a **knowledge base**, containing rules derived from human experts, and an **inference engine**, responsible for applying those rules to specific problems. One of the most successful examples was XCON, used by Digital Equipment Corporation to configure complex computer systems. Its success demonstrated that AI could deliver tangible business value and triggered an “AI boom” in corporate environments.

Despite their success, expert systems suffered from significant drawbacks. They were notoriously *brittle*, often failing completely when confronted with situations outside their narrow scope. They also faced the *knowledge acquisition bottleneck*, as extracting and encoding expert knowledge into formal rules proved costly, slow, and error-prone.

3.6 The Second AI Winter and the Shift to Agents (Late 1980s – Mid 1990s)

By the late 1980s, the limitations of expert systems—combined with high development and maintenance costs—led to another collapse in commercial enthusiasm, known as the *Second AI Winter*.

This period of reduced expectations nevertheless produced important theoretical progress. Researchers began to move away from top-down, rule-based representations of *knowledge* and toward bottom-up models of *behavior*. In robotics, figures such as Rodney Brooks advocated for building intelligence from layers of simple, reactive behaviors that interacted directly with the physical world. At the same time, the concept of the *intelligent agent* was formalized: an autonomous, goal-directed entity that perceives its environment and acts to maximize performance. This agent-centric view provided a unifying framework that emphasized what systems *do*, rather than what they explicitly *know*.

3.7 The Modern Renaissance: Big Data and Deep Learning (Present)

The current renaissance in AI is driven by the convergence of the agent paradigm with two transformative developments: the explosion of data and the revival of neural networks.

The growth of the internet and digital infrastructure has produced massive datasets, often referred to as *big data*, which enable learning at a scale previously impossible. At the same time, advances in hardware—particularly graphics processing units (GPUs)—have made it feasible to train large, multi-layer neural networks. This combination has fueled the rise of *deep learning*, dramatically improving performance in perception, language understanding, and decision-making tasks.

Modern AI agents represent the culmination of these historical trends. They are grounded in the rational agent framework of perception and action, but their internal mechanisms are increasingly powered by deep learning models trained on vast datasets. As a result, today's systems come closer than ever to the long-standing vision of intelligent, autonomous machines operating in complex real-world environments.

4 Foundations of Artificial Intelligence Agents

4.1 Defining the "Agent"

Before a comprehensive discussion of *Artificial Intelligence* agents is possible, it is first necessary to establish a clear and foundational understanding of what an "agent" is in the context of computer science. The term itself is broad, but in this domain, it has a precise and important meaning that forms the basis for all subsequent concepts.

In its most general sense, an **agent** is anything that can be viewed as perceiving its environment through **sensors** and acting upon that environment through **actuators** (Russell & Norvig, 2020). This is the standard, widely accepted definition in the field. It is a simple yet powerful concept that provides a framework for analyzing autonomous systems.

The agent exists in a continuous interaction loop with its surroundings. It gathers information using its sensors (a "percept") and then chooses an action to perform using its actuators. The agent itself is the core component that processes percepts and decides on actions.

Figure 4.1. Agent–environment interaction loop
[Environment] → [Sensors] → [Agent] → [Actuators] → [Environment]

4.2 From Agents to *Rational Agents*

The definition provided above is intentionally broad. However, the goal of AI is to create systems that act *well*, or intelligently. This distinction leads us to the crucial concept of the **rational agent**.

In the context of AI, "rationality" is defined purely in terms of performance. An agent is considered rational if it acts in a way that is expected to maximize its success, given the information it has (Russell & Norvig, 2020). To make this concrete, we define a **performance measure**, which is an objective criterion that evaluates the "success" of an agent's behavior (e.g., cleanliness for a vacuum, safety for a car).

Therefore, a **rational agent** is defined as an agent that selects an action that is expected to maximize its performance measure, given the evidence provided by its percept sequence and its built-in knowledge. Rationality is not omniscience; it is making the best possible decision based on available information.

4.3 The PEAS Framework

To move from abstract concepts to practical design, a more formal method of specification is required. The PEAS framework (Performance Measure, Environment, Actuators, Sensors) is a standard methodology used to define an agent's task and context *before* considering internal implementation (Russell & Norvig, 2020).

To illustrate this concept, Table 4.1 provides several examples of common AI agents defined using the PEAS framework.

Table 4.1. Examples of agent definitions using the PEAS Framework

Autonomous Vacuum

- **Performance Measure:** Cleanliness, battery, time, noise
- **Environment:** Room, obstacles, floor surfaces, dock
- **Actuators:** Wheels, brushes, vacuum, speaker
- **Sensors:** Camera, IR sensors, bump, dirt sensor

Self-Driving Car

- **Performance Measure:** Safety, speed, legality, comfort
- **Environment:** Roads, traffic, pedestrians, weather
- **Actuators:** Steering, accelerator, brakes, signals
- **Sensors:** Cameras, LiDAR, GPS, sonar, Radar

Email Spam Filter

- **Performance Measure:** Accuracy (spam vs. ham detection)
 - **Environment:** User's email inbox, server
 - **Actuators:** Mark spam, move, delete
 - **Sensors:** Email content, sender info, metadata
-

4.3.1 P — Performance Measures

Within the PEAS framework, the specification of **performance measures** establishes the normative basis upon which an agent's behaviour is judged, compared, and ultimately justified. This component is not a mere reporting convention appended after implementation; rather, it constitutes the formal articulation of what it means for the agent to “do well” in its task context. In modelling terms, performance measures define the objective function, explicitly or implicitly, against which the agent's policies, plans, or learned behaviours are optimized. Consequently, the quality of the overall agent model is constrained by the quality of this specification: if the performance criterion is ambiguous, incomplete, or misaligned with the intended task, then even technically sophisticated agents may exhibit systematically undesirable behaviour while still scoring highly on the chosen metric.

A rigorous performance specification must balance **clarity**, **operationalizability**, and **faithfulness to stakeholder intent**. Clarity demands that the measure be unambiguous and sufficiently formal to support reproducible evaluation. Operationalizability requires that the measure be computable from available observations, whether online during agent operation or offline in retrospective analysis, without introducing hidden degrees of freedom that permit post hoc reinterpretation. Faithfulness to stakeholder intent addresses a deeper modelling challenge: in many domains, the desired outcome is multi-faceted and cannot be reduced to a single scalar without loss. Here, explicit acknowledgement of trade-offs is essential, often motivating composite measures, constraint-based formulations, or multi-objective evaluation regimes. The PEAS framework encourages the modeller to make these choices visible rather than burying them within implementation details, thereby strengthening both scientific accountability and the interpretability of results.

Performance measures also play a central role in shaping the agent's **inductive biases** and failure modes, because the agent tends to exploit any systematic gap between the metric and the underlying desideratum. This phenomenon, commonly discussed in terms of metric gaming, reward hacking, or objective mis-specification,

underscores the importance of designing measures that are robust to superficial optimization. In practice, this robustness may require combining outcome-oriented metrics (e.g., task completion rate) with process- or constraint-oriented metrics (e.g., safety violations, resource budgets, fairness constraints), and reporting them separately rather than collapsing them into a single score. Moreover, careful attention must be paid to how evaluation is conducted across contexts: measures that appear valid under narrow laboratory conditions may fail to capture real-world costs, rare events, distributional shifts, or interactions with human users. A well-constructed performance specification therefore anticipates and resists the most likely forms of misgeneralization and unintended optimization.

Finally, performance measures create the evaluative bridge between conceptual modelling and empirical methodology. They determine which baselines are meaningful, what constitutes a fair comparison between alternative agent designs, and how progress can be measured over time. In a thesis, this is particularly consequential: claims about the efficacy, generality, or safety of a modelling approach are only as credible as the evaluation criteria that ground them. Accordingly, the “P” in PEAS should be treated as the first commitment in agent modelling, not because it is the only relevant component, but because it fixes the reference point against which all other design decisions (environmental assumptions, sensing fidelity, and actuation capabilities) must be interpreted. By specifying performance measures early and rigorously, the modeller ensures that subsequent modelling choices remain oriented toward the intended function of the agent, rather than toward incidental artefacts of a convenient or easily optimized metric.

4.3.2 E — Environment

In the PEAS framework, the **environment** denotes the external context within which an agent is situated and to which its behaviour must be responsive. The environment is not merely a backdrop for action but the structured source of uncertainty, constraint, and consequence that gives the agent’s decision-making problem its specific character. Accordingly, modelling the environment entails more than naming a domain (e.g., “a warehouse” or “a financial market”); it requires specifying those properties of the world that determine what information is available, how the world evolves, and how the agent’s actions alter future states. This component therefore functions as the principal determinant of problem complexity, because it shapes the space of admissible strategies and the conditions under which performance measures are meaningfully evaluated.

A central methodological task in environment specification is to characterize the **dynamics and informational structure** that govern agent–world interaction. Key distinctions include whether the environment is deterministic or stochastic, episodic or sequential, static or dynamic, discrete or continuous, and fully observable or partially observable. These distinctions are not taxonomic formalities; they guide the

selection and justification of modelling assumptions. For example, partial observability motivates belief-state reasoning or memory mechanisms, while non-stationarity raises questions about continual learning, robustness, and adaptation. Similarly, environments with delayed effects and long horizons compel attention to credit assignment and temporal abstraction. By requiring these properties to be made explicit, the PEAS framework prevents the modeller from inadvertently importing idealized assumptions (such as perfect observability or stationary dynamics) that may inflate apparent competence while undermining external validity.

Environment modelling also forces explicit engagement with **constraints** and **boundaries**, which are often the decisive factors in real-world agent performance. Constraints include physical limitations, resource budgets (time, energy, computation), legal and institutional rules, and safety-critical invariants that must not be violated. Boundaries determine what is treated as endogenous (within the model) versus exogenous (outside it), thereby fixing what counts as an “environmental disturbance” rather than an “agent failure.” Inadequate boundary-setting can yield misleading interpretations of results: an agent may appear ineffective simply because critical environmental features were omitted, or conversely appear effective because the environment was simplified in ways that remove the very difficulties the task is meant to address. A rigorous environment specification therefore serves as a safeguard against both over- and under-modeling, ensuring that empirical claims are anchored to a defensible representation of the task context.

Finally, the environment component provides the conceptual link between abstract agent modelling and the requirements of deployment and generalization. Agents rarely operate in a single, fixed world; they encounter families of environments with varying parameters, novel configurations, or adversarial perturbations. Explicitly modelling the environment as a distribution rather than a single scenario enables the thesis to distinguish performance in nominal conditions from robustness under shift. Moreover, environment specification clarifies the locus of generalization: whether the agent must generalize across states within one environment, across multiple related environments, or across qualitatively different regimes. In this way, “E” functions as the principal vehicle for establishing the scope and limits of the modelling enterprise, ensuring that later claims about competence, reliability, and transfer are framed with respect to precisely articulated environmental assumptions.

4.3.3 A — Actuators

In the PEAS framework, **actuators** specify the channels through which an agent can exert influence on its environment, thereby defining the agent’s action space in both form and scope. This component is foundational because it fixes what it means, operationally, for an agent to “act”: regardless of how sophisticated its internal reasoning may be, the agent’s competence is ultimately realized through the actions it is capable of taking. Actuator specification therefore functions as a concrete

constraint on achievable performance, delimiting which strategies are even expressible and which outcomes are reachable. From a modelling standpoint, it is insufficient to assume an abstract set of actions; one must instead characterize the available interventions in a way that is consistent with the environment and commensurate with the performance measures established in the “P” component.

A rigorous treatment of actuators requires attention to the **structure, granularity, and fidelity** of action. Actions may be discrete (selecting among symbolic operators) or continuous (issuing real-valued control signals), instantaneous or temporally extended, low-level (motor torques, cursor movements) or high-level (issuing commands, generating plans), and may vary in cost, risk, or reversibility. These characteristics matter because they shape the computational and statistical demands placed on the agent: fine-grained continuous control tends to require different modelling and learning approaches than sparse, high-level decision-making under constraints. Moreover, actuator fidelity (capturing limits such as noise, latency, saturation, or execution failure) often determines whether a model is externally valid. Ignoring such constraints can lead to “paper agents” that perform well in simulation yet fail when actuator imperfections are introduced, revealing that earlier success depended on unrealistic action capabilities rather than robust decision-making.

Actuator modelling is also inseparable from questions of **safety and governance**, since action channels are precisely the mechanisms by which an agent can cause harm or violate constraints. In many applied settings, the most important design decision is not how to maximize capability but how to bound it: limiting an agent’s actuators can reduce risk, improve interpretability, and enable effective oversight. This can take the form of action masking, permissioned tool use, rate limits, human-in-the-loop confirmation steps, or constrained optimization formulations that make certain interventions infeasible regardless of the agent’s preferences. Within the PEAS framework, such restrictions should be treated as first-class modelling commitments rather than after-the-fact “guardrails,” because they directly shape what kinds of policies the agent can implement and what failure modes are plausible.

Finally, actuators provide a crucial interface between abstract decision-making and measurable outcomes, and thus they mediate the alignment between the model’s internal objectives and the external world. The same performance measure can correspond to radically different agent designs depending on whether the actuators permit direct control, indirect influence, communicative action (e.g., recommendations or dialogue), or purely informational interventions (e.g., flagging anomalies). Explicit actuator specification clarifies these differences and prevents conflation between distinct forms of agency, such as an agent that “solves” a task by altering its environment versus one that solves it by persuading a human operator. In this sense, “A” stabilizes the interpretation of experimental results: it ensures that observed performance is attributed to the agent’s decision-making under defined action capabilities, rather than to implicit or uncontrolled degrees of freedom in how actions are represented and executed.

4.3.4 S — Sensors

In the PEAS framework, **sensors** specify the informational interface by which an agent acquires evidence about its environment, its own state, and the consequences of its actions. This component is decisive because it defines the epistemic limits of the agent: what it can know, when it can know it, and with what degree of reliability. In modelling terms, sensors determine the observation space and, by extension, the inferential burden placed on the agent's internal machinery. An agent cannot be evaluated as though it were omniscient if its sensory inputs are partial, delayed, noisy, or biased; conversely, an agent may appear unusually capable if the model implicitly grants it privileged access to latent variables unavailable in the intended deployment setting. Sensor specification thus functions as a guardrail against unrealistic assumptions and as a principled basis for interpreting both competence and failure.

A rigorous account of sensors must address the **scope, resolution, and quality** of observations. Scope concerns which aspects of the environment are observable at all; resolution concerns the granularity and modality of the signals (e.g., low-dimensional symbolic features versus high-dimensional raw data such as images, audio, or text); and quality concerns imperfections such as noise, missingness, calibration error, and systematic measurement bias. These properties materially affect the difficulty of the task: partial observability transforms decision-making into a problem of inference over hidden state; low temporal resolution can introduce aliasing and confound causal attribution; and biased measurements can induce persistent errors even in otherwise optimal policies. Accordingly, sensor modelling should not be treated as a passive catalogue of inputs but as a substantive commitment that shapes the agent's representational requirements, learning dynamics, and achievable performance.

Sensor specification also clarifies the relationship between **perception and decision-making**, preventing a common modelling conflation in which perceptual competence is assumed rather than demonstrated. When sensors deliver raw, high-dimensional signals, the agent must solve a perceptual problem (feature extraction, grounding, denoising, or semantic interpretation) before it can plan or act effectively. When sensors instead provide curated state variables, much of the perceptual difficulty is offloaded to the modeller, and the resulting agent model may better reflect control competence than end-to-end autonomy. A thesis that is explicit about this distinction can more precisely locate contributions: whether improvements arise from better perception, better policy optimization, better memory under partial observability, or better interaction protocols. In this way, “S” enables sharper scientific claims by tying performance to the informational conditions under which it was obtained.

Finally, sensors are central to questions of **robustness, validation, and responsible deployment**. Real-world sensing pipelines are often subject to distribution shift, adversarial manipulation, sensor drift, and unanticipated failure modes; these phenomena can dominate overall risk even when the underlying decision-making algorithm is sound. Treating sensors as first-class modelling elements encourages the inclusion of stress tests (e.g., noise injection, occlusion, missing-data regimes, domain shifts) that reveal whether an agent's competence is stable under plausible perturbations. Moreover, explicit sensor specification supports accountability by delineating what data the agent is permitted to use and what forms of monitoring are available for auditing its behaviour. Thus, in the PEAS framework, "S" is not merely the final letter of an acronym but the formal statement of the agent's epistemic boundary: an essential determinant of both what the agent can achieve and how confidently its performance can be interpreted.

4.4 The Importance of a PEAS-First Approach to the Modelling of AI Agents

A PEAS-first approach—an approach that foregrounds the specification of **performance measures**, the **environment**, **actuators**, and **sensors** prior to selecting algorithms or architectures—serves as a methodological anchor for the modelling of AI agents. In a thesis context, this stance is best understood as a disciplined inversion of a common tendency in contemporary practice: to begin with a preferred model class and subsequently retrofit a task description around it. By contrast, a PEAS-first posture forces the modeller to treat the agent as an embedded, goal-directed system whose competence is inseparable from the conditions of its operation. This reframing is not merely pedagogical; it establishes the minimal semantic commitments required to make claims about rationality, capability, and generalization meaningful. Consequently, it functions as a conceptual precondition for any rigorous analysis of agent behaviour, because it specifies what “success” is, where and how information is obtained, and by what channels the agent may intervene.

The primary contribution of a PEAS-first approach lies in its ability to prevent category errors during problem formulation, especially those arising from underspecified objectives. Performance measures, when articulated early, operate as a normative contract between modeller and system: they delimit the behavioural desiderata and render trade-offs explicit. This is especially important in settings where multiple metrics compete (e.g., accuracy versus latency, reward maximization versus constraint satisfaction, or user value versus safety criteria). Without such early commitments, modelling choices risk optimizing proxy objectives that are poorly aligned with the intended function of the agent, producing brittle systems that appear effective under narrow evaluation regimes yet fail under distributional shifts or adversarial conditions. A PEAS-first approach therefore promotes evaluability and comparability: it yields task definitions that are sufficiently precise to support ablation

studies, benchmarking, and principled error analysis, rather than retrospective narratives that rationalize observed performance.

Equally significant is the way PEAS-first modelling operationalizes the coupling between the agent and its environment. The environment component compels a structured account of the dynamics, uncertainties, and constraints within which an agent must act, while sensors and actuators concretize the informational and causal interfaces that mediate this interaction. This explicit interface specification has direct implications for what forms of reasoning and learning are feasible: partially observable environments, delayed consequences, or limited actuation fidelity impose qualitatively different demands than fully observable, instantaneous, and deterministic settings. In turn, a clear PEAS instantiation reduces ambiguity regarding whether a modelling challenge is fundamentally perceptual, inferential, planning-oriented, or control-theoretic. By treating perception and action as first-class commitments rather than implementation details, the approach discourages unrealistic assumptions (e.g., perfect state access or unconstrained control) that can inflate apparent competence and erode external validity.

A PEAS-first approach also improves the interpretability and governance of agent design decisions by making implicit modelling assumptions explicit and auditable. When sensors, actuators, and environment are defined upfront, one can more readily identify where failures originate: from insufficient sensing, limited action expressivity, mismatched environmental assumptions, or poorly chosen performance criteria. This traceability is essential not only for debugging but for responsible deployment, where stakeholders require intelligible accounts of system boundaries and risks. Moreover, early specification clarifies what constitutes permissible intervention and what data may legitimately inform decisions, thereby supporting alignment with legal, ethical, and domain-specific constraints. In this sense, PEAS-first modelling functions as a form of requirements engineering for agentic systems: it structures the problem space so that subsequent algorithmic choices are defensible responses to articulated constraints rather than opaque artifacts of convenience.

Finally, positioning PEAS as the starting point of modelling strengthens the cumulative coherence of a thesis that proceeds from conceptual foundations to technical instantiation. Having established the PEAS framework in the preceding subchapter and examined each component in isolation, the present argument emphasizes their methodological unity: PEAS is not merely a descriptive checklist but a generative blueprint for constructing agent models whose evaluation claims are meaningful, whose interfaces are well-defined, and whose failures are diagnosable. In doing so, a PEAS-first approach provides continuity between theoretical characterization and empirical implementation, ensuring that later chapters, whether they address learning algorithms, planning methods, or hybrid architectures, remain anchored to a stable and explicit task formalization. The result is a modelling pipeline that is both scientifically rigorous and practically robust, enabling conclusions about agent performance that generalize beyond a single experimental setup and remain interpretable under changing operational conditions.

5 Architectures of AI Agents

Having established the foundational concepts of agents, rationality, and environments, we turn to the internal structures, or **architectures**, that enable an agent to function. We will explore these architectures, moving from the simplest to the most complex.

5.1 Simple Reflex Agents

The **Simple Reflex Agent** is the most basic type. Its decision-making is straightforward: it bases its actions *only* on the **current percept**, ignoring history. Its behavior is analogous to a biological reflex. Its structure is based on **condition-action rules** (IF-THEN statements).

Conceptual structure (textual diagram):

Sensor percept → **Condition-action rule** → Action → Actuator

The primary advantage of this architecture is its simplicity and speed. The significant limitation is that it can only function effectively in **fully observable environments**.

5.2 Model-Based Reflex Agents

The **Model-Based Reflex Agent** addresses the problem of partial observability by maintaining an **internal state**. This state represents aspects of the environment the agent cannot currently perceive, built by remembering past percepts and using a **model** of how the world works.

Conceptual structure (textual diagram):

Percept → **Internal state and model** → **Condition-action rule** → Action

A common example is a robotic vacuum returning to its dock. It cannot always see the dock, but it remembers its location using an internal map (model).

5.3 Goal-Based Agents

The **Goal-Based Agent** is more advanced because it acts to achieve a specific **goal**, defined as a desirable state of the world. This architecture introduces **search** and **planning**. Instead of merely reacting, the agent considers future actions and evaluates sequences of actions that may lead to the goal.

Conceptual structure (textual diagram):

State → **Goal and planning** → Action sequence → Actuators

A standard example is a GPS navigation system, which plans an entire route from a current location to a destination goal.

5.4 Utility-Based Agents

The **Utility-Based Agent** is generally considered the most advanced rational agent architecture. It addresses the limitation of binary goals by using a **utility function** to measure how desirable a particular state is. The agent seeks to maximize **expected utility**, allowing it to handle uncertainty and conflicting objectives.

For example, a self-driving car may balance speed, safety, comfort, and fuel efficiency rather than simply reaching a destination as quickly as possible.

Conceptual structure (textual diagram):

State → **Utility function and planning** → Optimal action → Actuators

5.5 Summary of Architectures

The key distinctions between the four primary agent architectures are summarized below.

Table 5.1: Comparison of Primary AI Agent Architectures

Architecture	Basis for Action	Handles Partial Observability?	Handles Planning?	Handles Conflicting Goals?
--------------	------------------	--------------------------------	-------------------	----------------------------

Simple Reflex	Current percept	No	No	No
Model-Based Reflex	Internal state	Yes	No	No
Goal-Based	Current state + goal	Yes	Yes	No
Utility-Based	Current state + utility	Yes	Yes	Yes

6 In-Depth Case Studies of Agent Implementations

6.1 Introduction to Case Studies

The preceding chapters have established the theoretical foundations of Artificial Intelligence agents, defined their core components through the PEAS framework, classified the environments in which they operate, and detailed the primary architectures that govern their behavior. While these theoretical constructs are essential for a fundamental understanding of the field, the true power and complexity of AI agents are best understood through the examination of concrete, real-world implementations.

This chapter moves beyond the abstract to provide a rigorous, in-depth analysis of specific AI agent examples. By applying the theoretical tools developed in previous chapters to these case studies, we can gain a more nuanced appreciation for the practical challenges involved in designing autonomous systems. We will examine three distinct classes of agents: a chess-playing agent, an autonomous robotic vacuum cleaner, and a self-driving car.

6.2 Case Study 1: The Chess-Playing Agent

A chess-playing program is a quintessential example of an AI agent designed for a specific, well-defined task in a complex environment.

6.2.1 PEAS Analysis

- Performance Measure: Winning the game. Secondary measures include maintaining a material advantage or controlling key squares.
- Environment: The 8x8 chess board, the set of pieces, and the rules of movement. Crucially, it includes an opponent agent with conflicting goals.
- Actuators: Software functions to state a move in algebraic notation (e.g., “e2 to e4”).
- Sensors: A direct digital feed of the board configuration.

6.2.2 Environmental Classification

The chess environment is classified as:

- Fully Observable: The entire board state is visible at all times.
- Deterministic: The outcome of a move is completely determined by the current state and the action.
- Static: The board does not change while the agent is deliberating.
- Sequential: Every move has long-term consequences affecting future states.
- Multi-Agent (Competitive): The agent plays against a rational adversary.

6.2.3 Architectural Analysis

Given this classification, a chess agent must be a Goal-Based Agent. It uses sophisticated search algorithms (like minimax with alpha-beta pruning) to find a sequence of moves leading to the “checkmate” goal. The agent projects a vast tree of possible future moves and counter-moves, applying a heuristic evaluation function to estimate the value of future states, allowing it to select the optimal current move.

6.3 Case Study 2: The Autonomous Robotic Vacuum Cleaner

This case study moves to the physical, messy reality of a domestic environment, presenting challenges of uncertainty and incomplete information.

6.3.1 PEAS Analysis

- Performance Measure: Cleanliness of the floor, minimizing time, maximizing battery life, and avoiding getting stuck or damaging furniture.
- Environment: A home with various floor surfaces, static obstacles (walls, furniture), and dynamic obstacles (pets, people).

- Actuators: Wheel motors for movement, brush motors for cleaning, and a vacuum suction motor.
- Sensors: Bump sensors, cliff sensors, wall sensors, optical encoders (odometry), and dirt sensors.

6.3.2 Environmental Classification

The robotic vacuum environment is classified as:

- Partially Observable: Sensors only provide local information; the agent cannot see the entire house at once.
- Stochastic: Wheel slippage and sensor noise mean action outcomes are uncertain.
- Dynamic: The environment can change (e.g., people moving) while the robot works.
- Sequential: Current location depends on previous movements; cleaning requires a sequence of actions.
- Single-Agent: It is the only entity actively cleaning.

6.3.3 Architectural Analysis

A robust robotic vacuum must be a Model-Based and Goal-Based Agent.

1. Model-Based: To handle partial observability, it maintains an internal map (state) of the area explored and its location within it, often using techniques like SLAM (Simultaneous Localization and Mapping).
2. Goal-Based: It uses planning to achieve the goal of covering uncleaned areas efficiently and returning to the dock when the battery is low.

Simplified logic flow for a robotic vacuum agent:

Percept → Update Internal Map (Model) → Plan Path to Uncleaned Area (Goal) → Action

6.4 Case Study 3: The Self-Driving Car

The self-driving car is one of the most complex applications of AI agent technology, operating in a high-stakes, real-world environment.

6.4.1 PEAS Analysis

- Performance Measure: Hierarchical priorities: Safety (avoiding collisions) is paramount, followed by legality, efficiency, and passenger comfort.
- Environment: Public roads with diverse infrastructure, static objects, and scores of dynamic, independent agents (cars, pedestrians, cyclists), subject to varying weather and lighting.
- Actuators: Steering, throttle, and brakes.
- Sensors: A redundant array including multiple Cameras, LiDAR, Radar, and GPS/IMU.

6.4.2 Environmental Classification

The autonomous driving environment is the most challenging combination:

- Partially Observable: Occlusions prevent seeing the entire environment.
- Stochastic: Other agents behave unpredictably; sensors are noisy.
- Dynamic: The world changes rapidly and constantly.
- Sequential: Driving is a continuous sequence of dependent decisions.
- Multi-Agent: The agent shares the environment with countless other independent agents.

6.4.3 Architectural Analysis

Given the complexity and conflicting goals, a self-driving car must be a Utility-Based Agent. It uses a hierarchical architecture:

1. Perception: Fuses sensor data to create a coherent internal model of the world.
2. Prediction: Anticipates future behavior of other agents.
3. Planning (Utility Function): Generates multiple candidate trajectories and evaluates them using a complex utility function that heavily penalizes unsafe or illegal actions while rewarding efficiency and comfort. It selects the trajectory with the highest expected utility.
4. Control: Converts the chosen trajectory into actuator commands.

Simplified logic flow for an autonomous vehicle agent:

Sensor Fusion → Internal World Model → Trajectory Planning (Max Utility) → Control Commands

7 Applications and Societal Impact

The agent architectures and foundational principles described in the previous chapters are not merely theoretical constructs developed in academic isolation. They are the foundational blueprints for a wide array of systems that are currently deployed and are becoming increasingly integrated into the fabric of modern society. AI agents have moved from the laboratory into the consumer, commercial, and industrial sectors, often becoming so commonplace that they are no longer recognized by the general public as a form of artificial intelligence. This chapter will review some of the most prominent real-world applications of AI agent technology, outline the significant technical challenges that remain, and provide an in-depth examination of the profound ethical considerations that accompany their widespread adoption.

7.1 Real-World Applications

The deployment of AI agents is widespread across numerous sectors, validating the practical utility of the theoretical models discussed previously (Gartner, 2023).

Virtual Personal Assistants: Perhaps the most ubiquitous example of AI agents, systems like Apple's Siri, Amazon's Alexa, and the Google Assistant are now embedded in billions of devices worldwide (Kaplan, 2019). These software agents operate in a dynamic, multi-agent (human user), and partially observable environment. They utilize sensors (microphones) to perceive spoken language and actuators (speakers, screen displays) to respond. While often appearing as simple reflex agents in basic interactions, complex queries involve sophisticated model-based logic to maintain context and goal-based planning to execute tasks like booking reservations or sending messages.

Recommendation Engines: The algorithms powering platforms like Netflix, Amazon, and Spotify are powerful examples of utility-based agents operating at a massive scale. The “goal” of these systems is not merely to recommend a product, but to recommend the specific item that a user is most likely to value at that moment (Marr, 2020). They build complex internal models of user preferences based on vast histories of percepts (viewing habits, purchases, ratings) and use these models to maximize a utility function geared toward user engagement and retention.

Autonomous Systems and Robotics: In the physical domain, AI agents are most visibly represented by autonomous robotics. The industrial robotic arm used in manufacturing is a classic goal-based agent operating in a highly structured environment. The domestic autonomous vacuum cleaner is a model-based agent navigating a moderately complex environment. At the apex of current complexity is the self-driving car, a utility-based agent that must navigate highly dynamic, stochastic, and safety-critical environments.

7.2 Technical Challenges

Despite the successful deployment of these applications, the field of AI agent engineering is still nascent and faces significant technical hurdles that limit broader adoption in critical domains.

1. **Explainability (The “Black Box” Problem):** Many of the most powerful modern agents, particularly those utilizing deep neural networks for their internal models and utility functions, suffer from a severe lack of transparency. This is widely known as the “black box” problem. An agent may make a decision that

is statistically optimal according to its utility function, but human observers—including the system’s own developers—may be entirely unable to understand the specific chain of reasoning that led to that decision. The opaque nature of these complex mathematical models makes debugging difficult and poses severe challenges for establishing trust in safety-critical systems like medical diagnosis or autonomous driving (Shneiderman, 2020).

2. Safety and Reliability: For an agent to be trusted with autonomous operation in the real world, it must be demonstrably safe and reliable. This presents the formidable “alignment problem”: the challenge of ensuring that an agent’s programmed goals and utility functions are perfectly aligned with complex, nuanced human values and intentions. An agent pursuing a poorly defined goal with perfect rationality can lead to disastrous, unintended consequences. Furthermore, agents based on generative models may exhibit non-deterministic behavior or “hallucinations,” producing plausible-sounding but factually incorrect outputs, undermining reliability.
3. Data Quality and Dependency: Modern AI agents are not explicitly programmed with knowledge; they are “trained” on vast datasets. The performance and behavior of an agent are therefore fundamentally limited by the quality, scope, and veracity of the data it is fed. This leads to the classic computer science adage: “garbage in, garbage out.” If training data is incomplete, outdated, or contains hidden biases, the resulting agent will inherently reflect and amplify those flaws in its operational behavior (Broussard, 2018).

7.3 Ethical Considerations in Autonomous Systems

Beyond the significant technical hurdles, the rise of autonomous AI agents introduces a new and complex landscape of profound ethical questions. As these systems become more integrated into critical decision-making processes in finance, healthcare, criminal justice, and daily life, it is imperative to consider the moral frameworks that should govern their design and use (Floridi, 2019). These considerations are not merely academic exercises; they have tangible, real-world consequences for individuals and society as a whole. This section explores four of the most critical ethical challenges: bias and fairness, privacy and surveillance, job displacement, and accountability.

7.3.1 Bias, Fairness, and Algorithmic Discrimination

A primary ethical concern surrounding the deployment of AI agents is the potential for

these systems to learn, perpetuate, and even amplify existing human biases. Agents are not created in a social vacuum; they are trained on data generated by human society, which is inherently fraught with historical prejudices and systemic inequalities. If the training data reflects these biases, a rational agent, in its pursuit of maximizing its defined performance measure, will learn these patterns as “correct” representations of the world and apply them in its decision-making processes (Broussard, 2018).

This phenomenon, known as algorithmic bias, leads to unfair or discriminatory outcomes that can have serious negative impacts on marginalized communities. This is not necessarily the result of malicious intent on the part of developers, but rather a consequence of agents optimizing for utility based on flawed or unrepresentative data.

Consider the hypothetical scenario of a utility-based agent designed to streamline hiring processes for a large corporation. The agent’s performance measure is to maximize the rate of successful hires who stay with the company for more than two years. To build its internal model of a “good candidate,” the agent is trained on tens of thousands of historical resumes submitted to the company over the past three decades, along with data on which applicants were hired and their subsequent job performance.

If the historical data reflects a long-standing societal trend where leadership positions were predominantly held by men, the agent may identify statistical correlations between male-associated traits (e.g., specific gendered language in resumes, membership in certain organizations, or even names) and the concept of a “successful hire.” Consequently, in its effort to maximize its utility function, the agent may systematically downgrade highly qualified female candidates, not because it has a concept of sexism, but because its data-driven model suggests they have a lower probability of success based on historical patterns. The agent is acting rationally based on its inputs, but the outcome is deeply unethical and discriminatory.

Similar issues arise in other critical domains. In criminal justice, agents used to predict recidivism risk have been shown to exhibit racial bias due to being trained on historical arrest data that reflects systemic over-policing of certain communities. In finance, goal-based agents used for loan approval may learn to use zip codes as proxies for race or socioeconomic status, effectively engaging in a digital form of “redlining.”

The challenge of ensuring fairness in AI agents is immense. It requires moving beyond purely technical optimization to incorporate ethical constraints directly into agent architectures and utility functions. It also necessitates rigorous auditing of training data and model outputs to identify and mitigate disparate impacts on protected groups. Ensuring that autonomous agents serve all members of society equitably remains an unresolved challenge in the field.

7.3.2 Privacy, Consent, and the Surveillance State

The efficacy of many modern AI agents is directly proportional to the amount of data they can access and process. A model-based agent builds a more accurate internal state of the world the more it observes; a utility-based agent makes better decisions for a user the more it knows about that user's preferences, habits, and context. This reliance on vast quantities of personal data creates a fundamental tension between the convenience and utility offered by AI agents and the human right to privacy.

Consider the ubiquitous virtual personal assistant installed in a smart home. To function effectively as a rational agent that can respond instantly to voice commands, the device's microphone must be in a state of constant passive listening, processing ambient audio to detect its "wake word." While manufacturers assure users that audio is not recorded or transmitted until the wake word is detected, the very existence of an always-on sensor in intimate private spaces represents a significant shift in privacy norms.

Furthermore, to provide personalized assistance (such as managing calendars, suggesting routes based on traffic, or recommending products) the agent must accumulate and analyze a staggering amount of behavioral data. This includes search history, location data tracked via smartphones, communication logs, purchase history, and media consumption habits. This data is used to construct highly detailed digital profiles of individuals, often without their full comprehension of the scope or granularity of the data collection.

This raises critical ethical questions regarding informed consent. The traditional model of presenting a user with a lengthy, complex Terms of Service agreement and requiring a single click to "agree" is widely viewed as inadequate for the era of pervasive AI agents. Users often do not truly understand what data is being collected, how long it is being stored, who it is being shared with, or the potential downstream consequences of that data being fed into complex algorithmic models.

Moreover, the infrastructure built for consumer convenience can easily be repurposed for surveillance. The network of sensors, cameras, and microphones connected to AI agents in smart homes, smart cities, and workplaces creates an unprecedented architecture for monitoring human behavior. This capability could be exploited by authoritarian governments to suppress dissent or by corporations to monitor employees with invasive granularity. The risk is the gradual normalization of a surveillance state where autonomous agents are the primary tools of observation and control, fundamentally eroding individual autonomy and privacy.

7.3.3 Job Displacement and Economic Disruption

The economic and societal impact of widespread automation through AI agents is another major area of ethical concern. The fundamental purpose of creating autonomous agents is to automate tasks previously performed by humans. While this automation has the potential to significantly increase productivity, efficiency, and economic output, it also carries the substantial risk of displacing human workers on a massive scale and exacerbating economic inequality.

Unlike previous waves of technological automation that primarily affected manual or repetitive labor, AI agents are increasingly capable of performing cognitive, non-repetitive tasks. Agents can now draft legal documents, analyze financial markets, diagnose medical conditions, write software code, and generate creative content. This means that a much broader segment of the workforce, including highly educated white-collar professionals, is now vulnerable to displacement.

Consider the potential impact of autonomous vehicle agents on the transportation and logistics industry. Truck driving is one of the most common occupations in many countries, providing stable, middle-class incomes for millions of individuals without college degrees. The widespread deployment of utility-based autonomous trucking fleets, which can operate 24 hours a day without rest and prioritize fuel efficiency perfectly, would offer immense economic benefits to shipping companies and consumers through lower costs.

However, the human cost would be profound. The rapid displacement of millions of truck drivers would devastate local economies dependent on their income and create a massive social challenge of retraining a large workforce for new occupations. History suggests that such transitions are rarely smooth and that the economic gains from automation are rarely distributed equitably. There is a significant risk that the economic benefits of AI agents will accrue primarily to the owners of the capital (the technology companies and large corporations that deploy the agents), while the costs of adjustment fall disproportionately on displaced workers, widening the wealth gap and leading to social instability. The ethical challenge here is not just about the technology itself, but about the social and economic policies required to manage the transition and ensure that the benefits of AI-driven abundance are shared broadly across society.

7.3.4 The Accountability Gap

Finally, the increasing autonomy of AI agents creates a profound legal and ethical dilemma known as the “accountability gap” or the “responsibility vacuum.” In

traditional engineering, if a machine or tool fails and causes harm, it is generally clear who is responsible: the operator who misused it, or the manufacturer who designed it with a flaw. However, autonomous AI agents break this traditional chain of causality.

When an agent is designed to learn from its environment and make independent decisions without human intervention in real-time, determining responsibility for harmful outcomes becomes incredibly difficult. If a complex, utility-based autonomous vehicle is involved in a fatal accident, who is held accountable?

Is it the developers who wrote the initial code? They might argue that they could not have foreseen the specific, complex scenario that led to the accident and that the agent's behavior evolved based on its training data. Is it the company that provided the training data? They might argue that the data was industry standard. Is it the owner of the vehicle? They were likely just a passenger, relying on the system's promise of autonomy. It cannot be the agent itself, as an AI software program has no legal standing, no assets to pay damages, and no moral capacity to be punished or held responsible.

This problem is severely compounded by the “black box” issue of explainability discussed in the technical challenges section. If human investigators cannot determine why the agent made the decision that led to the accident—if the decision was the result of millions of opaque calculations within a deep neural network—it becomes nearly impossible to assign negligence or liability under current legal frameworks.

This accountability gap is a critical ethical flaw. A fundamental principle of justice is that if harm occurs, there must be a mechanism for redress and accountability. If society cannot clearly define who is responsible when autonomous agents cause harm, it will be impossible to build the necessary public trust to deploy these systems in high-stakes domains like healthcare, transportation, and law enforcement. Resolving this gap requires not just technological solutions for better explainability, but likely significant changes to legal frameworks regarding liability for autonomous systems.

8 Conclusion and Future Outlook

8.1 Summary of Thesis

This thesis has provided a comprehensive and foundational exploration of the subject of Artificial Intelligence Agents. The objective was to synthesize and review the established, core concepts that define this field, situate them historically, and analyze their real-world implications.

Chapter 1 introduced the topic, defined key terms, and set the context of the current AI era. Chapter 2 provided a review of the foundational academic, industrial, and critical literature. Chapter 3 offered necessary historical context, tracing the evolution of the field from early cybernetics through various paradigms to the modern data-driven agent approach.

Chapter 4 established the conceptual groundwork, defining the “rational agent,” the PEAS framework, and classifying environments. Chapter 5 examined the primary architectures of AI agents, progressing from simple reflex models to complex utility-based systems.

Chapter 6 applied these theoretical concepts through in-depth case studies of a chess agent, a robotic vacuum, and an autonomous car. Finally, Chapter 7 discussed real-world applications, technical challenges, and provided an expanded examination of pressing ethical considerations including bias, privacy, jobs, and accountability.

8.2 The Future of AI Agents

The field of AI agents is dynamic and rapidly evolving. The future points toward two major trends: increased complexity in individual agents and increased interconnectedness between agents.

One significant frontier is Multi-Agent Systems (MAS), studying the interactions of many autonomous agents within a single environment. This is vital for applications like managing autonomous logistics fleets, optimizing city-wide traffic, or modeling complex economic systems.

Another key trend is the integration of agents with the Internet of Things (IoT). Future agents will holistically manage complex systems like smart homes or industrial plants by processing and acting upon data from vast networks of connected devices. Furthermore, advances in deep learning and reinforcement learning will continue to make agents more adaptive, capable of learning their own models and utility functions directly from raw interaction with the world, reducing the reliance on human pre-programming.

8.3 Final Word

This thesis has charted the landscape of AI Agents, from basic theoretical definitions to complex real-world implementations. The potential for these autonomous systems to bring about positive transformation in efficiency, safety, and convenience across all sectors of society is undeniable.

However, as discussed extensively in Chapter 7, this immense potential is inextricably linked with significant risks. The formidable technical challenges of ensuring safety, reliability, and explainability, combined with the profound ethical questions of algorithmic bias, the erosion of privacy, economic disruption, and the accountability gap, are central to the technology's future.

In conclusion, the field of AI agents is a dual-use technology of unprecedented power. As we move forward, the responsible development of these systems will require not only continued technical innovation but also careful, thoughtful, and proactive governance. Further interdisciplinary research and a robust societal dialogue are essential to navigate the complex future that autonomous AI agents will help create.

References (mandatory item – NBR14724, item 4.2.3.1; NBR 6023)

RUSSELL, Stuart J.; NORVIG, Peter. *Artificial intelligence: a modern approach*. 4th ed. [S. I.]: Pearson, 2020. ([dblp](#))

BROUSSARD, Meredith. *Artificial unintelligence: how computers misunderstand the world*. [S. I.]: The MIT Press, 2018. ([MIT Press](#))

FLORIDI, Luciano; COWLS, Josh. A unified framework of five principles for AI in society. *Harvard Data Science Review*, [S. I.], v. 1, n. 1, 2019. Available at: <https://doi.org/10.1162/99608f92.8cd550d1>. Accessed on: Dec. 17, 2025. ([Directory of Open Access Journals](#))

GARTNER, Inc. *Gartner top 10 strategic technology trends for 2024: AI trust, risk and security management (AI TRiSM)*. [S. I.]: Gartner, 2023. Available at: <https://www.gartner.com/en/articles/gartner-top-10-strategic-technology-trends-for-2024>. Accessed on: Dec. 17, 2025. ([Gartner](#))

KAPLAN, Andreas; HAENLEIN, Michael. Siri, Siri, in my hand: who's the fairest in the land? On the interpretations, illustrations, and implications of artificial intelligence. *Business Horizons*, [S. I.], v. 62, n. 1, p. 15-25, 2019. Available at: <https://doi.org/10.1016/j.bushor.2018.08.004>. Accessed on: Dec. 17, 2025. ([ScienceDirect](#))

MARR, Bernard. The 10 best examples of how AI is already used in our everyday life. *Forbes*, [S. I.], Dec. 16, 2019. Available at: <https://www.forbes.com/sites/bernardmarr/2019/12/16/the-10-best-examples-of-how-ai-is-already-used-in-our-everyday-life/>. Accessed on: Dec. 17, 2025. ([Forbes](#))

SHNEIDERMAN, Ben. Human-centered artificial intelligence: reliable, safe & trustworthy. *International Journal of Human–Computer Interaction*, [S. I.], v. 36, n. 6, p. 495-504, 2020. Available at: <https://doi.org/10.1080/10447318.2020.1741118>. Accessed on: Dec. 17, 2025. ([Taylor & Francis Online](#))

TURING, Alan M. Computing machinery and intelligence. *Mind*, [S. I.], v. 59, n. 236, p. 433-460, Oct. 1950. Available at: <https://doi.org/10.1093/mind/LIX.236.433>. Accessed on: Dec. 17, 2025. ([OUP Academic](#))

WOOLDRIDGE, Michael. *An introduction to multiagent systems*. 2nd ed. [S. I.]: John Wiley & Sons, 2009. ([Wiley](#))